



# NoSQL: Graph Database



Presenter: Longbin Lai



# Background



# From SQL Database to NoSQL Database

---

- What is SQL Database?
  - Simply: What you have studied prior to this course, a.k.a. relational database
  - Features:
    - The data is well-structured (schema)
    - Querying using SQL
- What is NoSQL Database?
  - No standard definition, “Not only SQL”
  - Features:
    - The data can be stored in any format according to applications: Text file, Key/Value, XML, ...
    - Querying is often ad-hoc, but more often borrow from SQL

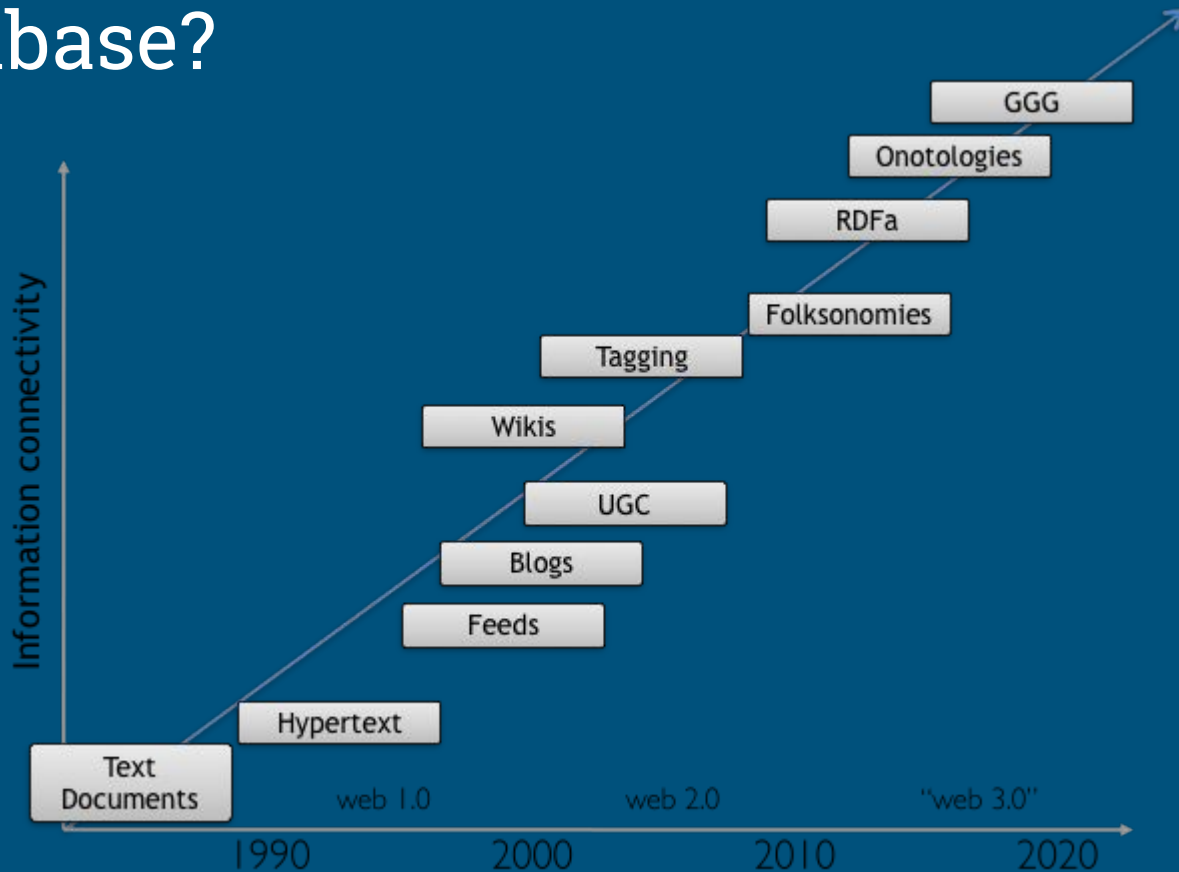
# Why NoSQL Databases?

- Data is getting bigger
- How much you are using in you postgresql?

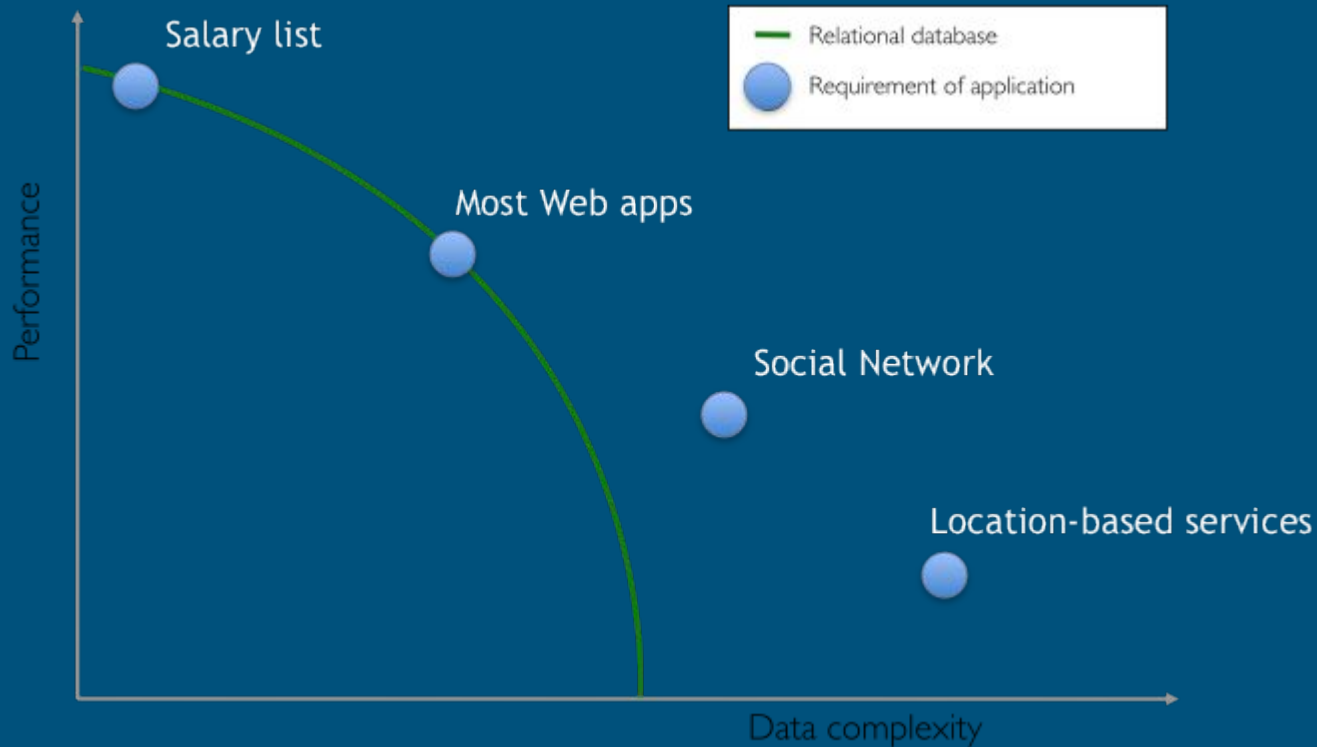


# Why NoSQL Database?

- Data is more connected



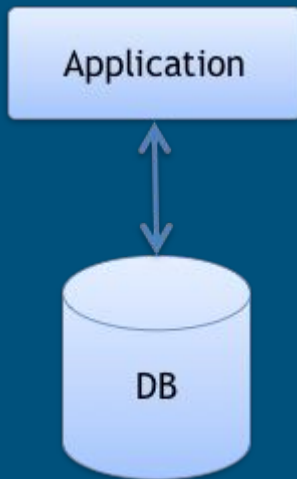
# RDBMS performance



# Why NoSQL Database

---

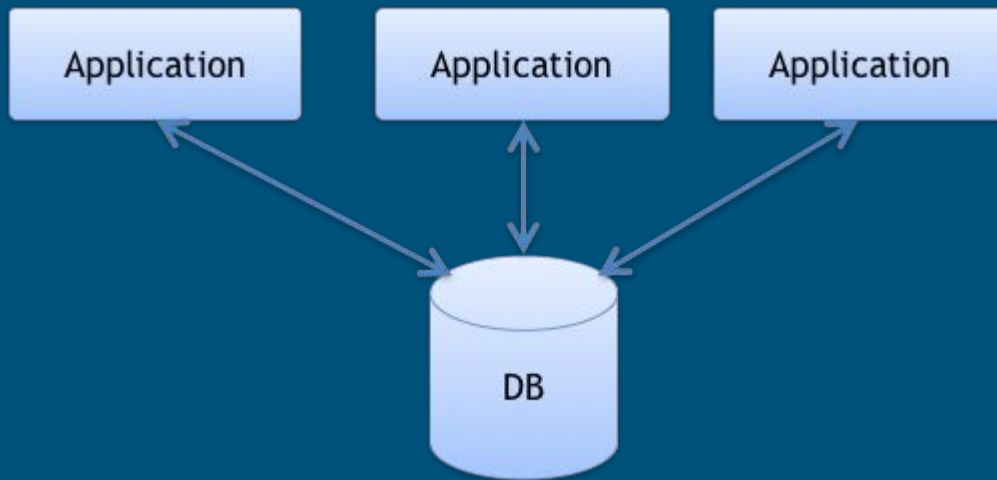
- 1980': Single Application



# Why NoSQL Database

---

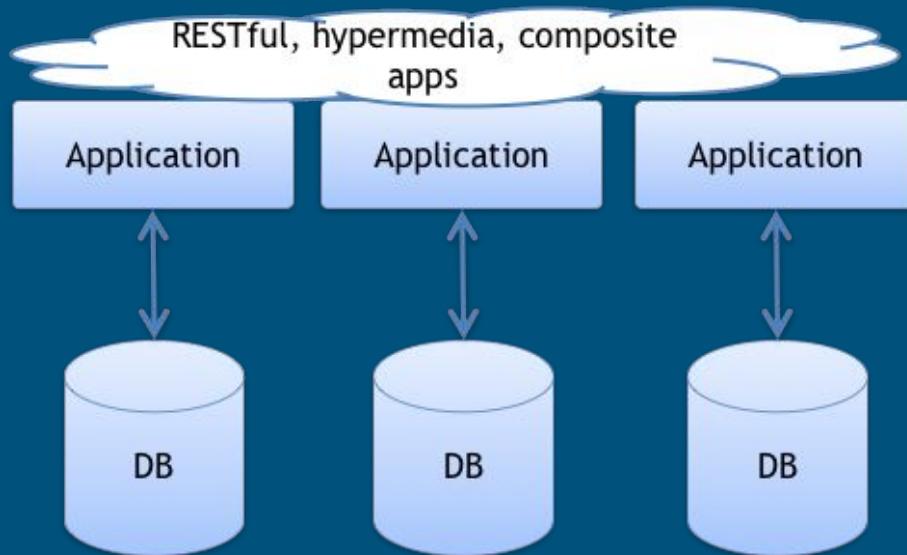
- 1990's: Integration Database





# Why NoSQL Database

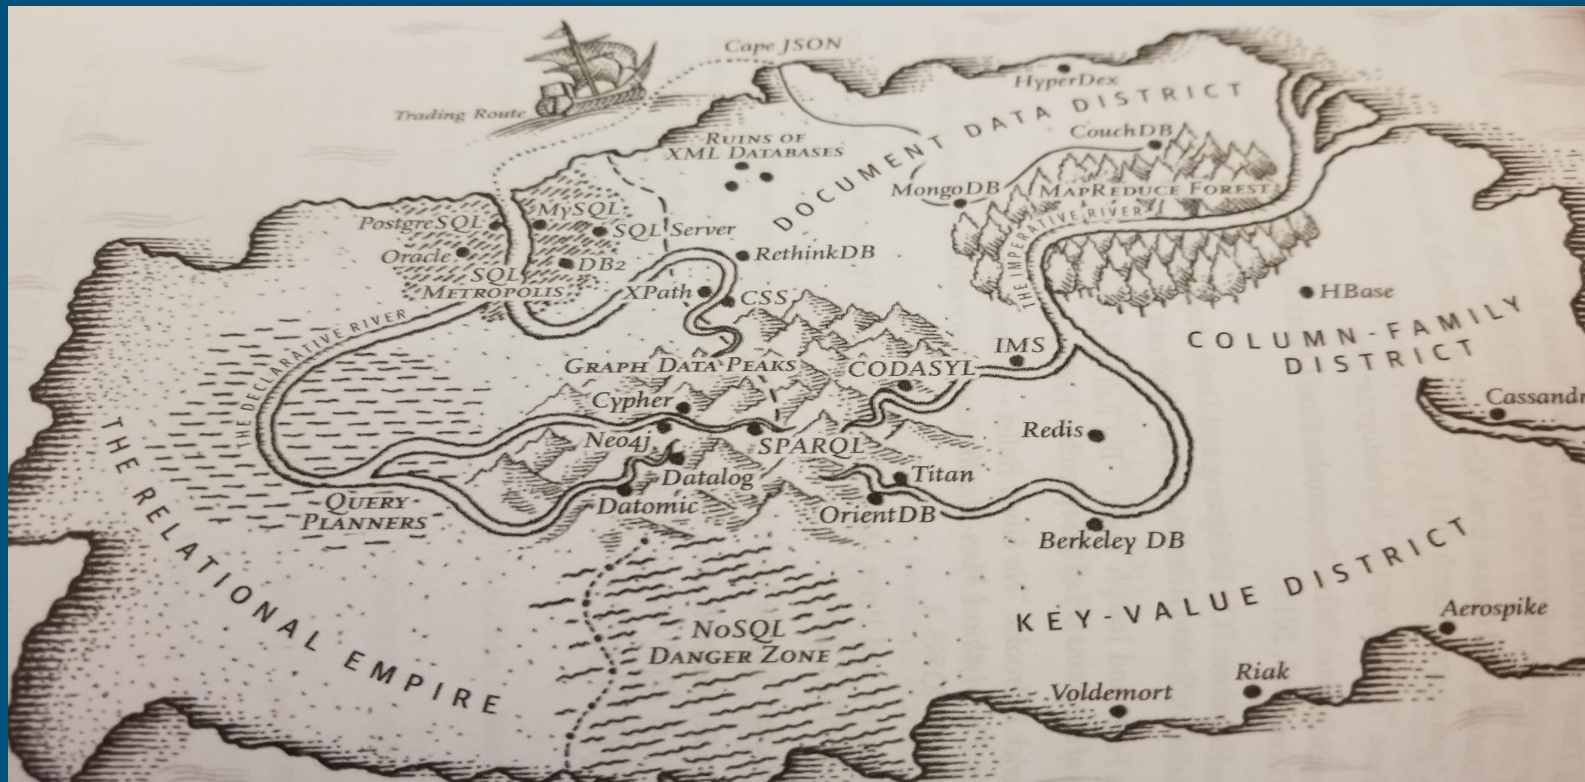
- 2000' ~ Service-based



We start to have multi-sourced data, which means impossibility of maintain one common structure

# The DB World Map

Book: Design Data-Intensive Applications



# NoSQL

Not Only SQL



# NoSQL Vendors

---



# Key-Value Stores

---

- Come from a research article by Amazon (Dynamo)
  - Global Distributed Hash Table (Key-Value Stores)
- Popular Vendors
  - Redis (Open Sourced)
  - Amazon's DynamoDB (the inventor)
  - Microsoft Azure Cosmos DB

# Everything in Key-Value

---

- Friendship of Facebook
  - Relational Database:
    - People (Id, Name, ...)
    - Friend (Pid1, Pid2, Time)
  - Key-Value Store:
    - Profile KV: <Id; Name>
    - Friendship KV: <Id; Set{(Id1, Time1), (Id2, Time2), ...}>
- When we want to get all friends
  - Relational Database: Join People with Friend (costly)
  - Key-Value Store: Get directly from the Friendship KV ( $O(1)$ )

# Key-Value Store

---

- Why
  - Simple Data Model: Hash Table is mature data structure
  - Good Scalability: Small System Cost, via good look-up locality and caching
- Why not
  - Poor to complex (interconnected) data
    - many KV pairs needed to be maintained for each data
    - hash-table-like structure tends to performance poorly for large data

# Column Family

---

- Origin from Google's BigTable
- Main Idea
  - Each table tends to have many attributes (thousands ~ millions)
  - In most applications (analytics) we are only interested in a few (10s)
  - Traditional row-based
    - Store the each record in a sequential file
    - Even just read one attribute, we should read the whole record
  - Column-based
    - Store the data by putting the same attribute in a sequential file
    - Faster access to a few attributes



# Column Family

---

- Google's BigTable
  - Drives MapReduce
  - Apache Hadoop, Hadoop File System (HDFS), HBase
  - Apache Cassandra

# Column Family

---

- Why?
  - Optimized for data analytics
  - Semi-Structured Data: Each column can define its own schema
  - Big Data
- Why not?
  - Nightmare for interconnected data
    - Suppose I have an Friendship table (pID1, pID2, ...) stored based on column
    - Want to find how many steps from me to “Donald Trump”
    - Sequentially scan the “pID1” and “pID2” column again and again

# Others

---

- Document-based DB
  - XML
  - MongoDB (one of most popular DBs)
  - CouchDB
- Specific-purpose
  - RDF: 3-tuple (obj1, <action>, obj2), e.g. (John, add\_friend, Emily)
  - To resolve the problem of maintaining interconnected data.
  - But it does not give the ultimate solution

# Graph Database

---

- Data Model
  - Nodes (Vertices) -> Entities
  - Edges -> Relations
  - Are we going to learn ER model again?
- Main Vendors
  - Neo4j
  - JanusGraph
  - OrientDB
  - Gremlin

# Graph Database

---

- Why
  - Natural fitness for data with complex connections
  - Powerful data model, as general as relational model
  - Flexible query: based on many graph algorithms
- Why not
  - Not easy to scale, because of **poor locality**
  - Most algorithms are computationally intensive


# Graph is not just everywhere, it is \$\$\$


---

Recommendations 

Social Computing  

Geospatial (Google map) 

Internet of things 

Web Analytics 

Bioinformatics .. 

**Market Cap: US\$2 Trillion**

# Graph Database

---

# What is a Graph Data Structure

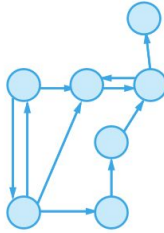
- Node / Vertex (id, name, age) -> Entity
- Edge / Link / Arc (srcId, dstId, timestamp) -> Relation



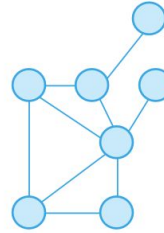


# Graph Types (Edge Types)

- Directed Graph
- Undirected Graph



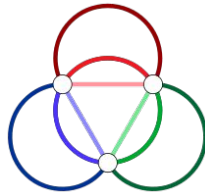
Directed



Undirected

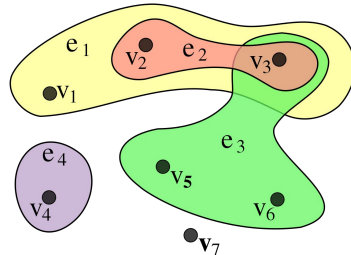
Edge: Two nodes  
Two nodes: at most one edge

- Multi Graph



Edge: Two nodes  
Two nodes:  $\geq$  one edge

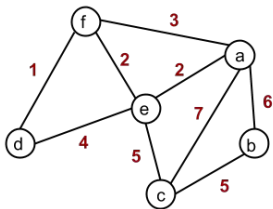
- HyperGraph



Edge:  $\geq$  two nodes  
Two nodes:  $\geq$  one edge

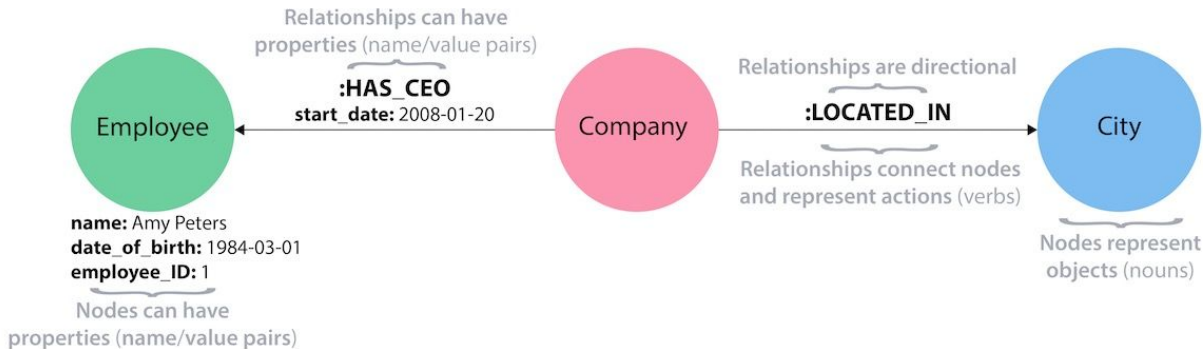
# Graph Types (Attributes)

- Weighted Graph



Every edge has a number as its **weight**

- **Property Graph**

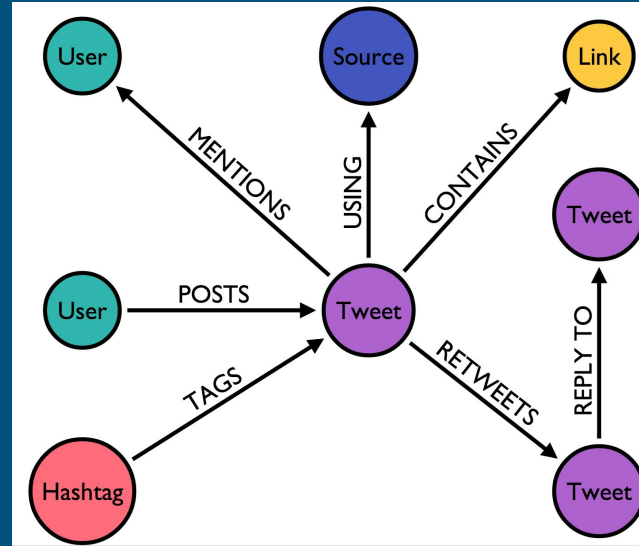
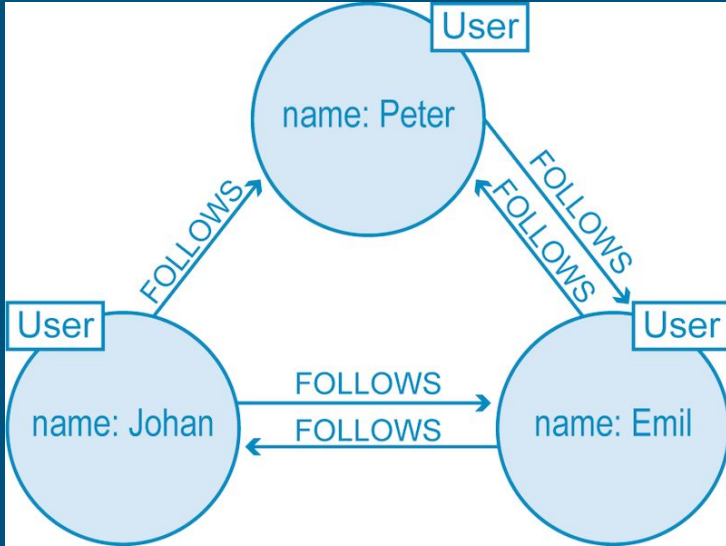


# Properties

---

- Analogous to Attributes in relational table
- Both Nodes and Edges can have properties
- Properties are key-value pairs
  - e.g. "name": John, "prof": musician, "join\_day": 09/04/2019
- Properties are more flexible than Attributes
  - Primitive type: String, Integer, Char, Boolean, ...
  - Array: int[ ], String[ ]
  - Set: int{ }, String{ } // Set can not contain duplicate elements

# Twitter Graph



“following”/“followed by” indicates directions

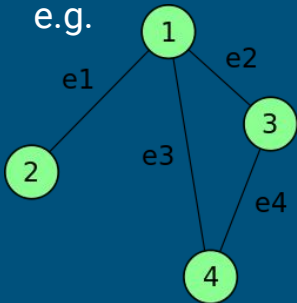
# Facebook Graph - undirected

**"friend" is a mutual relationship, which indicates undirection**

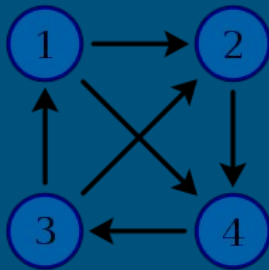
# Graph database

- A database based on an explicit graph structure
- Every node maintains not only its properties, but its adjacent (neighbor) nodes

- Adjacent nodes are nodes that I connect with
- e.g.



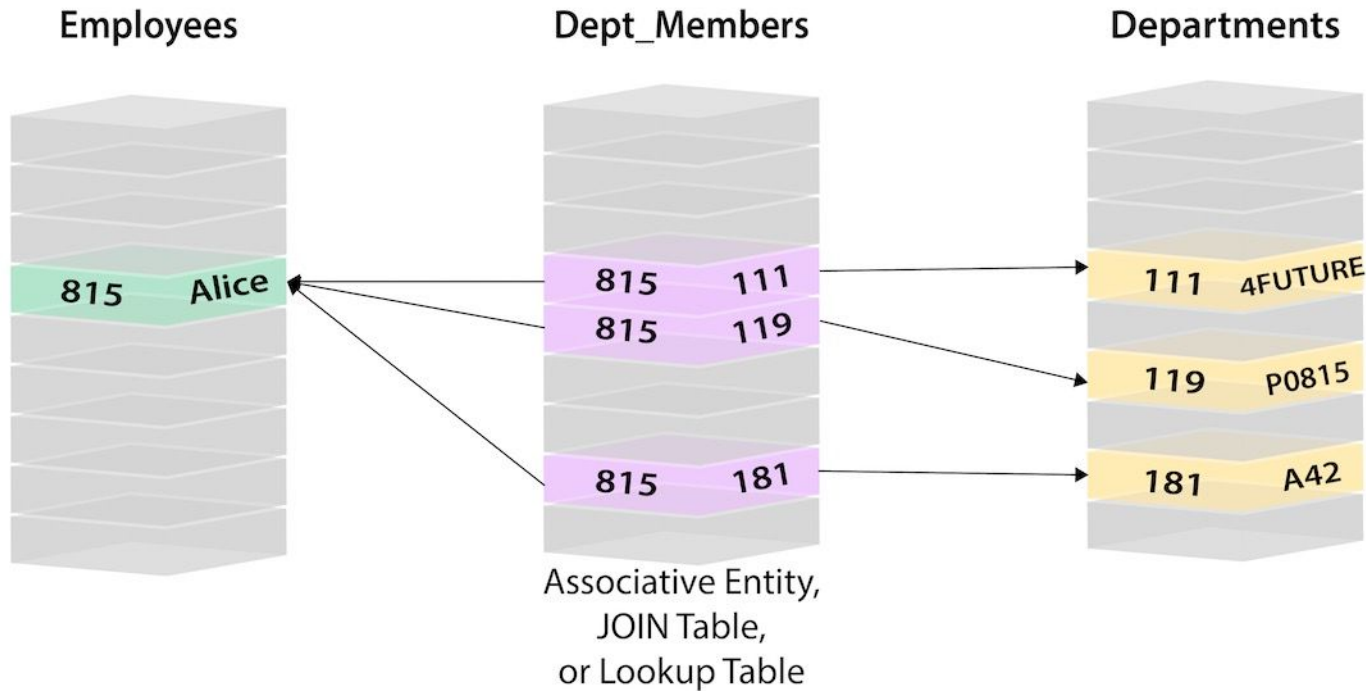
1: {2, 3}  
2: {1}  
3: {1, 4}  
4: {1, 3}



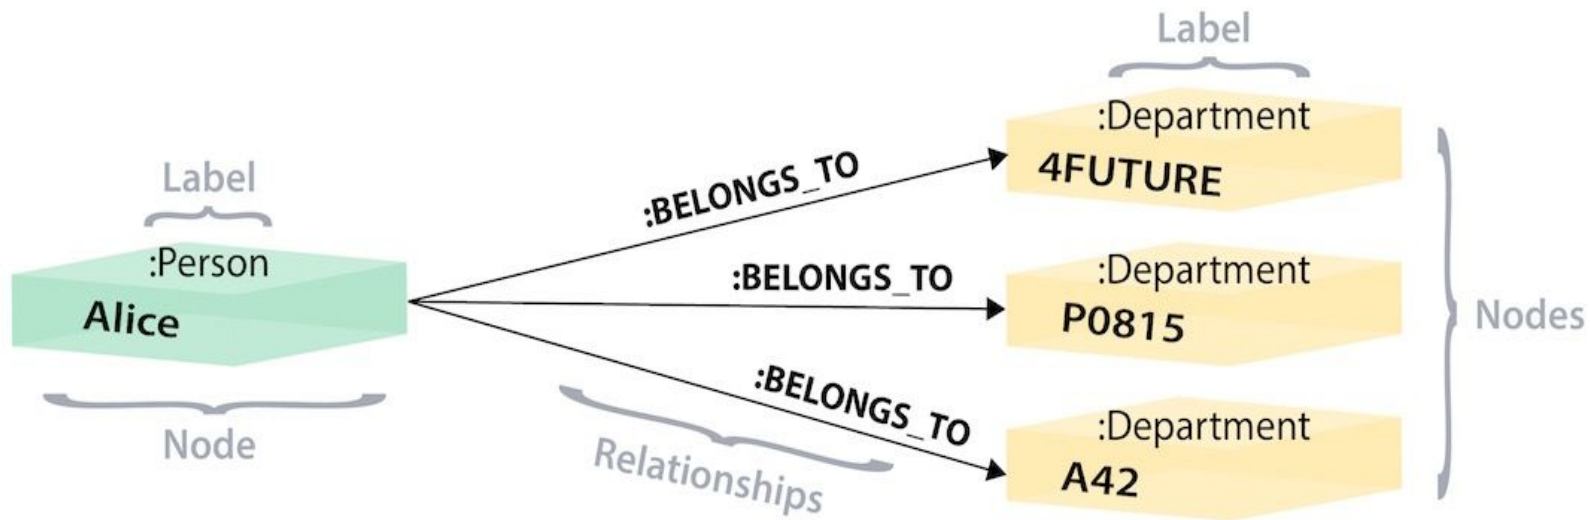
Directed Graph:  
In neighbors and out neighbors

1: Out {2, 4}, In {3}  
2: Out {4}, In {1, 3}  
3: Out {2}, In {4}  
4: Out {3}, In {1, 2}

# Relational Databases



# Graph Database





# Graph DB for Connection Analytics

---

- Find out among my friends' friends, who are also my friends?
- Solving with relational database

- People (Id, name)
- Friend (srcId, dstId, ...)
- 
- **SELECT** Distinct(p2.name)
- **FROM** People p1, People p2, Friend f1, Friend f2, Friend f3
- **WHERE** f1.dstId = f2.srcId **AND** # Get friend
- f2.dstId = f3.dstId **AND** # Get friend's friends
- f1.srcId = f3.srcId **AND** # Check also my friend
- f1.srcId = p1.id **AND**
- p1.name = "MY\_NAME" # check it is me
- f3.dstId = p2.id **AND** # get friend's name
-

# Graph DB for Connection Analytics

---

- Find out among my friends' friends, who are also my friends?
- Solving with graph database
  - `Node (srcId, name, nbrs(friends): [nbr1, nbr2, nbr3, ...])`
  - 
  - `Set {}` # Initialize an empty set
  - **GET** `my_node` from Graph DB # Get the node from Id
  - **FOR EACH** `nbr` **IN** `my_node.nbrs`: # Outer: Check all my friends
  - **FOR EACH** `nbr_nbr` **IN** `nbr.nbrs`: # Inner: Check my friend's friends
  - **IF** `nbr_nbr` **IS IN** `my_node.nbrs`: # Check also my friend
  - `Set.insert(nbr_nbr)`
  - **RETURN** `Set`

# Graph DB for Connection Analytics

---

- Find out the shortest path from me to Donald Trump in the Twitter network
  - A Path: A sequence of nodes where there is an edge in between
    - e.g. A -> B -> C -> D -> E (path of length 4)
  - A shortest path: For every such path, the one with shortest length
    - e.g. A -> B -> C -> D -> E -> F
    - A -> B -> C -> F **Shorter one**
- 
- Use everyday
  - Navigation (Road)
  - Routing (Network)
  -

# Graph DB for Connection Analytics

- Find out the shortest path from me to Donald Trump
- Solving with Relational DB

- STEP 1:
- `SELECT COUNT(*)`
- `FROM friend f1`
- `WHERE f1.srcId = myId AND f1.dstId = TrumpId`
- 
- `IF COUNT = 0 -> STEP 2:`
- `SELECT COUNT(*)`
- `FROM friend f1, friend f2`
- `WHERE f1.srcId = myId AND`
- `f1.dstId = f2.srcId`
- `f2.dstId = TrumpId`

**Goes until COUNT != 0,  
When completes?**



# Graph DB for Connection Analytics

---

- Single-Source Shortest Path Problem
  - Given a source node of the graph, find the shortest path of all nodes to this source node
  - Nightmare to solve using relational DB
  - Well-defined graph algorithm
    - Breadth-First Search (BFS): When all edges have the same weight value
    - Dijkstra Algorithm: When all edges have some positive weight
    - Bellman-Ford Algorithm: General case
  - Assume Directed Graph G
    - Each node {id: ID, dist: FLOAT, out\_nbrs: [ nodes ], is\_visited: BOOLEAN }
    - Each edge { srcId: ID, dstId: ID, weight: FLOAT }

# Breadth-First Search (BFS)

- Queue

- First In First Out (FIFO)
- EnQueue(v): Put an element into the queue
- DeQueue(): Return and remove the top element in the queue
- e.g.

1. Q.EnQueue(v1), Q.EnQueue(v2), Q.EnQueue(v3)

2. Q.DeQueue -> v1

3.

4. Q.EnQueue(v4)

5.

6. Q.DeQueue -> v4



Q.DeQueue -> v2

Q.DeQueue -> v3

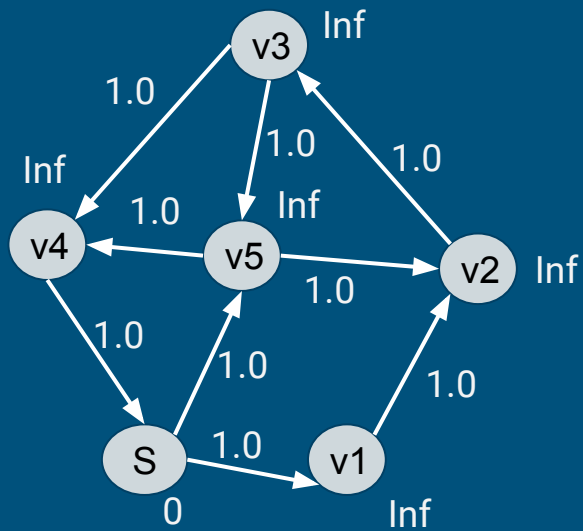
# Breadth-First Search

---

- Initialize every node's distance as **Infinity**
- Make the source node's distance as 0
- `Q.Enqueue(s)` // `s` is the source node
- Loop until `Q` is empty
  - `v = Q.DeQueue()`
  - for each `nbr` in `v.out_nbrs()`:
    - if `nbr.dist > v.dist + e(v, nbr).weight`:
      - `nbr.dist = v.dist + e(v, nbr).weight (=1)`
      - if not `nbr.is_visited`: `Q.Enqueue(nbr)`
  - Make `v` as visited: `v.is_visited = true`
- EndLoop

# Breadth-First Search

Initialization



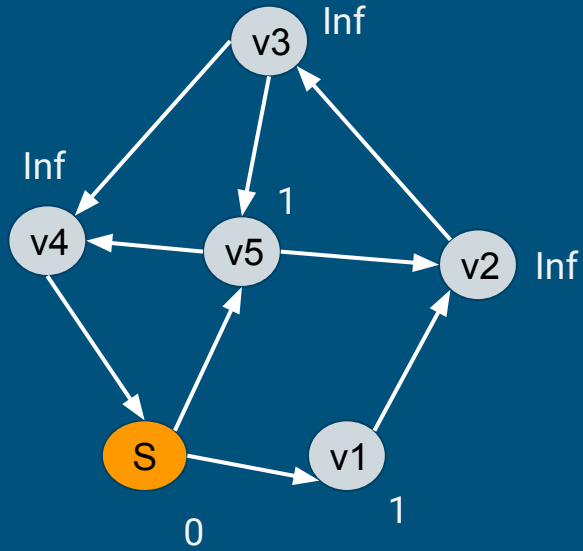
Q





# Breadth-First Search

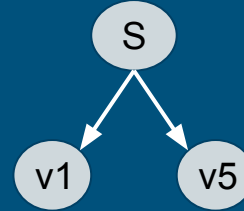
Loop1



Q



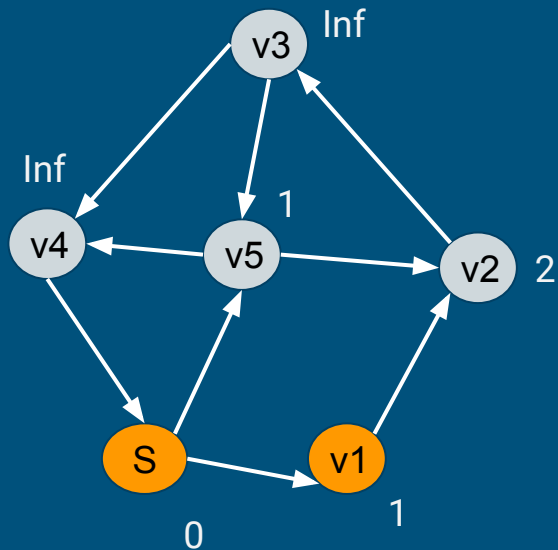
Q.dequeue



S.out\_nbrs()

# Breadth-First Search

Loop2



Q



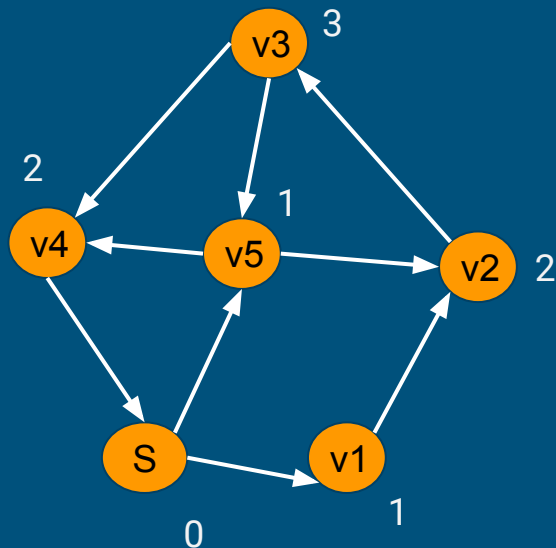
Q.dequeue



`v1.out_nbrs()`

# Breadth-First Search

Loop3~6



Loop3: DeQueue v5, update v4 {2}, EnQueue v4

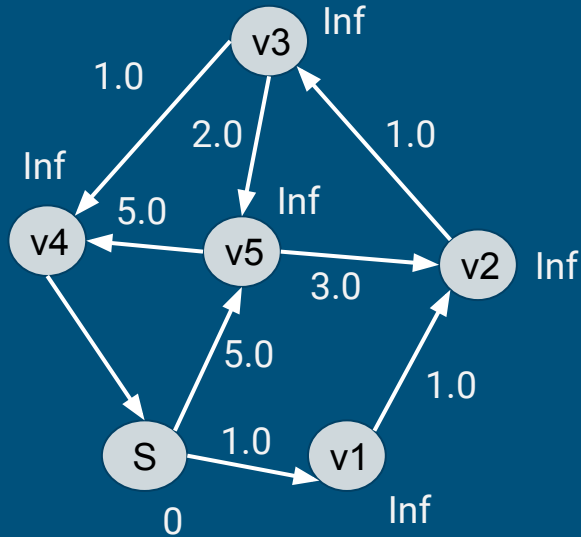
Loop4: DeQueue v2, update v3 {3}, EnQueue v3

Loop5: DeQueue v4, no ACT

Loop6: DeQueue v3, no ACT

# Breadth-First Search

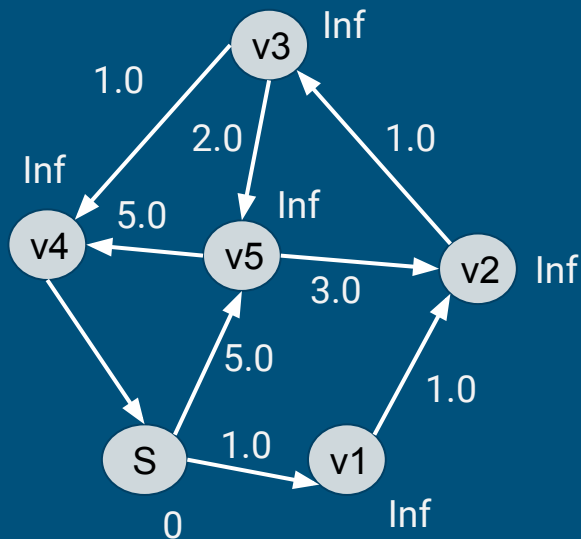
---



If Edge does not have the same weight value, BFS will compute **wrong** results?

Practice !!

# Dijkstra Algorithm



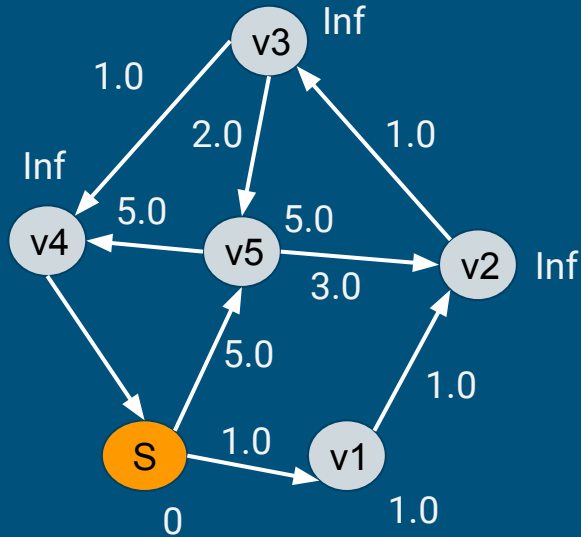
Dijkstra Algorithm: Queue -> PriorityQueue

Priority Queue

- Each element has a key
- DeQueue the element with the smallest key, rather than FIFO

# Dijkstra Algorithm

Loop1

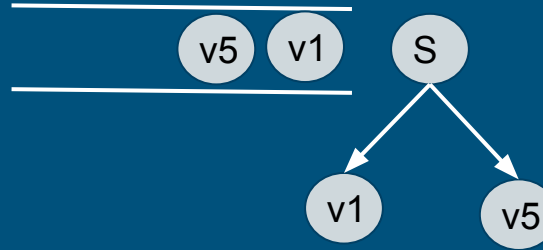


Dijkstra Algorithm: Queue -> PriorityQueue

Priority Queue

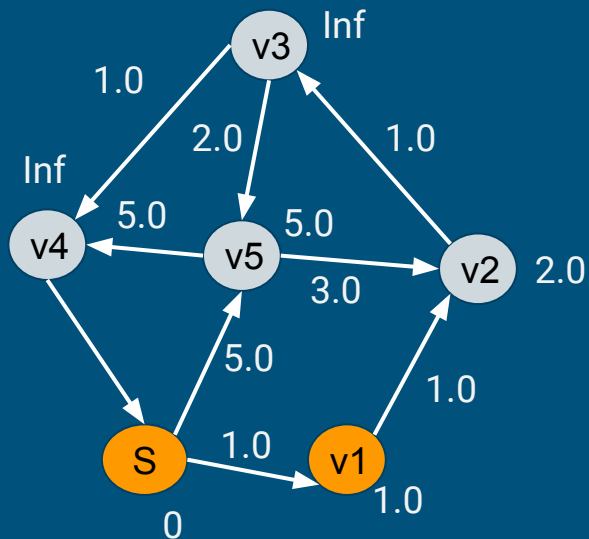
- Each element has a key
- Dequeue the element with the **smallest** key, rather than FIFO

PQ



# Dijkstra Algorithm

Loop2

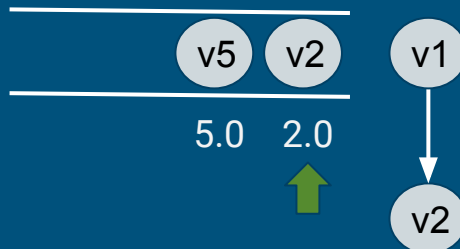


Dijkstra Algorithm: Queue -> PriorityQueue

Priority Queue

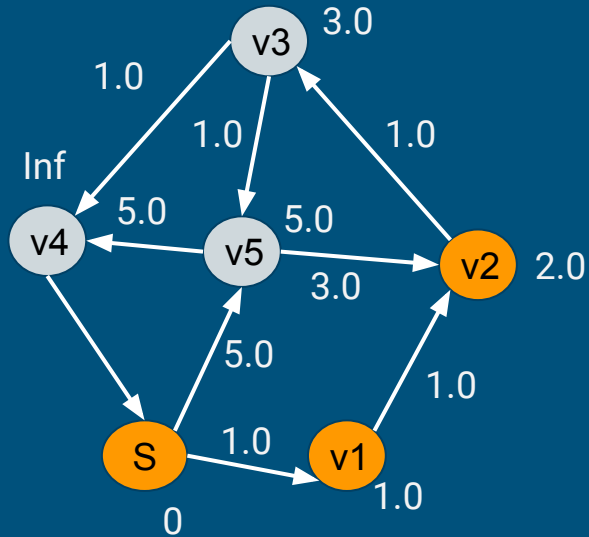
- Each element has a key
- Dequeue the element with the **smallest** key, rather than FIFO

PQ



# Dijkstra Algorithm

Loop3

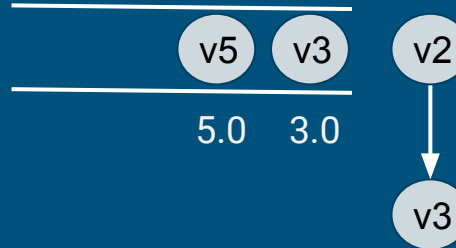


Dijkstra Algorithm: Queue -> PriorityQueue

Priority Queue

- Each element has a key
- Dequeue the element with the **smallest** key, rather than FIFO

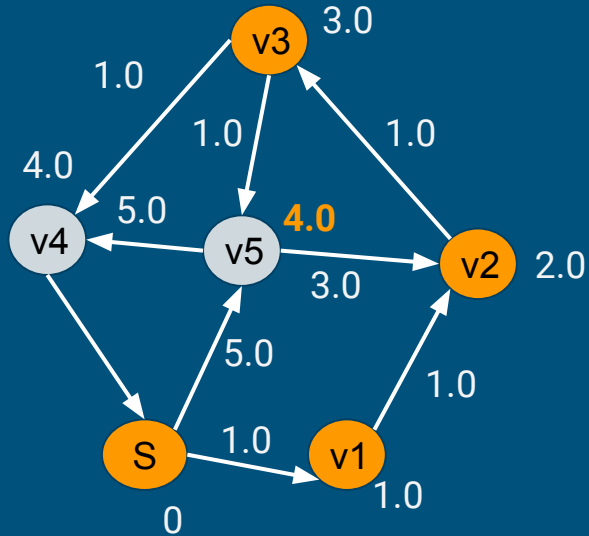
PQ





# Dijkstra Algorithm

Loop3

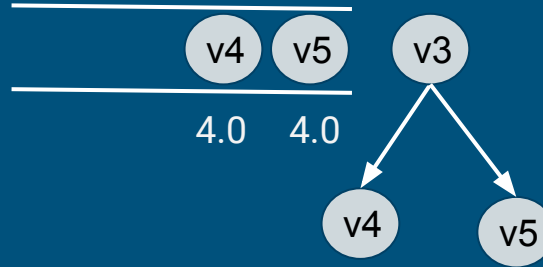


Dijkstra Algorithm: Queue -> PriorityQueue

Priority Queue

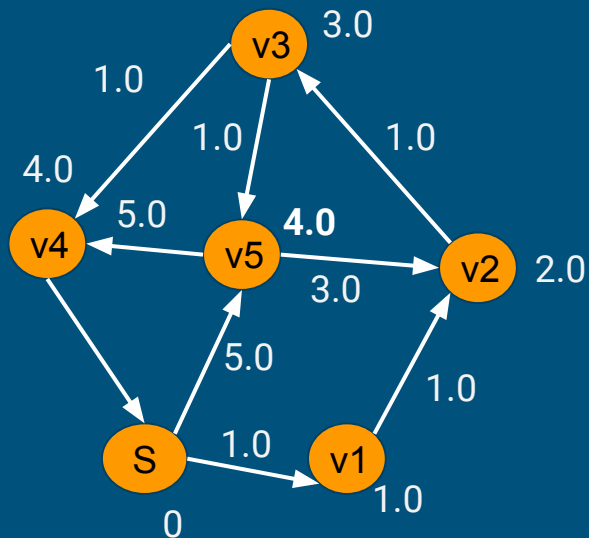
- Each element has a key
- Dequeue the element with the **smallest** key, rather than FIFO

PQ



# Dijkstra Algorithm

LoopEnd



Dijkstra Algorithm: Queue -> PriorityQueue

Priority Queue

- Each element has a key
- Dequeue the element with the **smallest** key, rather than FIFO

# Summarizations

---

# Summarizations

---

- Graph DB is NoSQL DB
- Graph DB's main structure
  - Node: Define entity, with properties (attributes), **nbrs (in\_nbrs, out\_nbrs)**
  - Edge: Define relationships (connections), directed or undirected
- Graph DB is specially designed for Connection Analytics
  - Friend's friends are friends (Triangle Listing)
  - Single-source shortest path (SSSP)

# Next Class

---

- Cypher Query Language in Neo4J
- Distributed Graph Processing
  - Pregel: Vertex-centric Computation Model
  - Some typical algorithms: Triangle Listing, PageRank, SSSP, etc.